

LLM-AS-A-JUDGE CHEATSHEET

AI Engineer Ready · Week 6 — RAG Optimization + Evaluation + LangSmith

Golden Rule -> a strong LLM grades another model's answer against a rubric -- fast and scalable, but not infallible

1. What LLM-as-a-Judge Is

	Detail
What it is	A strong LLM grades another model's answer against a rubric, like a teacher marking an exam.
Why use it	Far cheaper and faster than a human reviewing every answer -- scales to thousands of cases.
Watch out for	Judges can favor longer, more confident-sounding answers. Spot-check against real humans.

2. Pairwise vs. Absolute Scoring

Style	How It Works	When It's More Reliable
Absolute (1-5)	Judge scores a single answer on a fixed scale.	Quick to run and aggregate, but judges are inconsistent about where "4" starts.
Pairwise (A vs B)	Judge is shown two candidate answers and picks the better one.	Often more reliable -- comparing is an easier task for an LLM than assigning an absolute number.

Example: for "How do I reset my password?", Answer A ("Click Forgot Password, check email, set a new password") beats Answer B ("Reset through account settings if logged in, or via support if not") because it's more direct and matches how most users actually reset a password.

3. A Minimal Judge Prompt

```
$ SYSTEM: You are an impartial evaluator. Score the ANSWER for how well it
$ addresses the QUESTION, using only the given CONTEXT as ground truth.

$ Rubric:
$ 5 = fully correct, fully grounded, directly answers the question
$ 3 = partially correct or partially grounded
$ 1 = wrong, irrelevant, or unsupported by context

$ Return only the integer score.

$ QUESTION: {question}
$ CONTEXT: {context}
$ ANSWER: {answer}
```

4. Known Judge Biases

LLM-AS-A-JUDGE CHEATSHEET

AI Engineer Ready · Week 6 — RAG Optimization + Evaluation + LangSmith

Golden Rule -> a strong LLM grades another model's answer against a rubric -- fast and scalable, but not infallible

Bias	What Happens
Verbosity bias	Judges tend to rate longer answers higher, even when a shorter answer is equally correct.
Position bias	In pairwise comparisons, the judge can favor whichever answer appears first (or second) -- always randomize the order.
Self-preference bias	A model judging its own output family tends to score it more favorably than a competing model's output.

5. Best Practices

- Use a strong, capable model as the judge -- a weak judge produces noisy scores.
- Write a short, specific rubric; don't just ask "is this good?"
- Randomize answer order in pairwise comparisons to cancel out position bias.
- Spot-check a random sample against real human judgments regularly.
- Use the judge for the bulk of evaluation, not as the only signal you ever look at.

Pitfall

Judges can favor longer, more confident-sounding answers over shorter, correct ones. Never fully trust a judge score without periodically checking it against human agreement on a random sample.

Remember

LLM-as-a-Judge is a scale tool, not a truth oracle. Use it to grade thousands of cases cheaply, then let humans validate that the judge itself is trustworthy.

INTERVIEW PREP

Q: How do you evaluate a RAG system when you don't have ground-truth labels?

A: LLM-as-a-judge with a clear rubric is a standard, defensible answer -- a strong model grades faithfulness and relevance directly from the question, context, and answer, without needing a pre-written correct answer.

Q: Why might pairwise comparison be more reliable than a 1-5 score?

A: Judging "which of these two is better" is a simpler, more consistent task for an LLM than anchoring an absolute number on a scale, which tends to drift and disagree across runs.

Q: Name two known biases in LLM-as-a-judge and how you'd mitigate them.

A: Verbosity bias (favors longer answers) -- control for length or explicitly instruct the judge to ignore length. Position bias (favors first/second answer in a pairwise setup) -- randomize which answer is shown first.