

LangSmith Traceability, Recall, Precision & GenAI Evaluation Metrics

Q1. What is the primary purpose of traceability in LangSmith?

- A. To increase model training speed
- B. To track and analyze the execution flow of LLM applications
- C. To compress embeddings
- D. To reduce API costs

Q2. In LangSmith, a trace typically represents:

- A. A database backup
- B. A network connection
- C. A record of an application's execution and interactions
- D. A model checkpoint

Q3. Recall measures:

- A. How many retrieved results are relevant
- B. How many relevant results were successfully retrieved
- C. The speed of retrieval
- D. The size of the knowledge base

Q4. A retrieval system found 8 relevant documents out of 10 total relevant documents available. What is its Recall?

- A. 80%
- B. 20%
- C. 100%
- D. 50%

Q5. Precision measures:

- A. The percentage of retrieved items that are relevant
- B. The percentage of all relevant items retrieved
- C. The latency of the system
- D. The token count of responses

Q6. A search system retrieves 20 documents, but only 15 are relevant. What is the Precision?

- A. 50%
- B. 75%
- C. 25%
- D. 90%

Q7. Which evaluation metric is commonly used to determine whether an LLM response is factually correct?

- A. Correctness
- B. Temperature
- C. Token Usage
- D. Context Length

Q8. In a RAG application, low Recall usually indicates that:

- A. Too many irrelevant documents are retrieved
- B. Relevant information is being missed during retrieval
- C. The model is generating faster responses
- D. The embedding size is too large

Q9. Which LangSmith feature helps developers identify where an LLM application failed during execution?

- A. Prompt Templates
- B. Traces and Runs Visualization
- C. Vector Databases
- D. Fine-Tuning

Q10. Why are evaluation metrics important in Generative AI systems?

- A. They increase GPU memory
- B. They help measure and improve model quality and reliability
- C. They automatically train the model
- D. They eliminate hallucinations completely